

Lab 1 Solution: Linear Regression

```
def predict(x, w):
    """
    Compute the prediction of a linear model.
    Inputs:
    x: np.ndarray input data of shape [num_samples, num_feat + 1]
    w: np.ndarray weights of shape [num_feat + 1, 1]
    Outputs:
    h: np.ndarray predictions of shape [num_samples, 1]
    """
    h = np.dot(x,w)
    return h
```

```
def compute_cost(x, y, w):
    """
    Inputs:
    x: np.ndarray input data of shape [num_samples, num_feat + 1]
    y: np.ndarray targets data of shape [num_samples, 1]
    w: np.ndarray weights of shape [num_feat + 1, 1]
    Outputs:
    mse: scalar.
    """
    m = x.shape[0]
    mse = np.sum((y-predict(x,w))**2) / (2*m)
    return mse
```

```
def gradient_descent(x, y, w, learning_rate, num_iters):
    """
    Inputs:
    x: np.ndarray input data of shape [num_samples, num_feat + 1]
    y: np.ndarray targets data of shape [num_samples, 1]
    w: np.ndarray weights of shape [num_feat + 1, 1]
    learning_rate: scalar, the learning rate.
    num_iters: int, the number of iterations.
    Outputs:
    j_hist: list of loss values of shape [num_iters]
    w_opt: [num_feat + 1, 1]
    w_hist: [num_feat + 1, num_iters + 1]
    """
    n_sample, n_feat = len(x), len(w) - 1
    j_hist = np.zeros([num_iters])
    w_hist = w
```

```

for i in range(num_iters):
    h = np.dot(x,w)
    dw = np.dot((h-y).T,x).T / n_sample
    w = w - learning_rate*dw
    w_hist = np.append(w_hist,w,1)
    j_hist[i] = compute_cost(x,y,w)
w_opt = w_hist[:,(np.argmin(j_hist)+1)].reshape(-1,1)
return j_hist, w_opt, w_hist

```

```

def compute_cost_multivariate(x, y, w):
    """

```

Inputs:

x: np.ndarray input data of shape [num_samples, num_feat + 1]

y: np.ndarray targets data of shape [num_samples, 1]

w: np.ndarray weights of shape [num_feat + 1, 1]

Outputs:

mse: scalar.

"""

```

    m = x.shape[0]
    res = predict(x,w) - y
    mse = np.dot(res.T, res)[0,0] / (2*m)
    return mse

```